torch.cuda.comm.broadcast_coalesced#

https://pytorch.org/docs/stable/generated/torch.cuda.comm.broadcast_coalesce

Description:

Broadcast a sequence of tensors to the specified GPUs. Small tensors are first coalesced into a buffer to reduce the number of synchronizations. tensors (sequence) – tensors to broadcast. Must be on the same device, either CPU or GPU. devices (Iterable[torch.device, str or int]) – an iterable of GPU devices, among which to broadcast. buffer_size (int) – maximum size of the buffer used for coalescing

Function Signature

```
torch.cuda.comm.broadcast_coalesced(tensors, devices,
buffer_size=10485760)[source]#

Parameters

Returns
```

Returns:

A tuple containing copies of tensor, placed on devices.

Detailed Documentation

Documentation

Broadcast a sequence of tensors to the specified GPUs.

Small tensors are first coalesced into a buffer to reduce the number of synchronizations.

tensors (sequence) – tensors to broadcast. Must be on the same device, either CPU or GPU.

devices (Iterable[torch.device, str or int]) – an iterable of GPU devices, among which to broadcast.

buffer_size (int) – maximum size of the buffer used for coalescing

A tuple containing copies of tensor, placed on devices.

Broadcast a sequence of tensors to the specified GPUs.

Small tensors are first coalesced into a buffer to reduce the number of synchronizations.

tensors (sequence) – tensors to broadcast. Must be on the same device, either CPU or GPU.

tensors (sequence) – tensors to broadcast. Must be on the same device, either CPU or GPU.

devices (Iterable[torch.device, str or int]) – an iterable of GPU devices, among which to broadcast.

devices (Iterable[torch.device, str or int]) – an iterable of GPU devices, among which to broadcast.

buffer_size (int) – maximum size of the buffer used for coalescing

buffer_size (int) – maximum size of the buffer used for coalescing

A tuple containing copies of tensor, placed on devices.
