

torch.cuda.comm.reduce_add_coalesce

https://pytorch.org/docs/stable/generated/torch.cuda.comm.reduce_add_coalesce

Description:

Sum tensors from multiple GPUs. Small tensors are first coalesced into a buffer to reduce the number of synchronizations. inputs (Iterable[Iterable[Tensor]]) – iterable of iterables that contain tensors from a single device. destination (int, optional) – a device on which the output will be placed (default: current device). buffer_size (int) – maximum size of the buffer used for coalescing

Function Signature

```
torch.cuda.comm.reduce_add_coalesced(inputs,  
destination=None, buffer_size=10485760)[source]#
```

Parameters

Returns

Returns:

A tuple of tensors containing an elementwise sum of each group of inputs, placed on the destination device.

Detailed Documentation

Documentation

Sum tensors from multiple GPUs.

Small tensors are first coalesced into a buffer to reduce the number of synchronizations.

`inputs (Iterable[Iterable[Tensor]])` – iterable of iterables that contain tensors from a single device.

`destination (int, optional)` – a device on which the output will be placed (default: current device).

`buffer_size (int)` – maximum size of the buffer used for coalescing

A tuple of tensors containing an elementwise sum of each group of inputs, placed on the destination device.