

SDPBackend#

<https://pytorch.org/docs/stable/generated/torch.nn.attention.SDPBackend.html>

Description:

An enum-like class that contains the different backends for scaled dot product attention. This backend class is designed to be used with the `sdpa_kernel` context manager. **ERROR**: An error occurred when trying to determine the backend. **MATH**: The math backend for scaled dot product attention. **FLASH_ATTENTION**: The flash attention backend for scaled dot product attention. **EFFICIENT_ATTENTION**: The efficient attention backend for scaled dot product attention.

Function Signature

```
class torch.nn.attention.SDPBackend#
```

The following Enums are available:

```
property name#
```

Notes & Warnings

WARNING

Warning This class is in beta and subject to change.

Detailed Documentation

Documentation

An enum-like class that contains the different backends for scaled dot product attention. This backend class is designed to be used with the `sdpa_kernel` context manager.

ERROR: An error occurred when trying to determine the backend.

MATH: The math backend for scaled dot product attention.

FLASH_ATTENTION: The flash attention backend for scaled dot product attention.

EFFICIENT_ATTENTION: The efficient attention backend for scaled dot product attention.

CUDNN_ATTENTION: The cuDNN backend for scaled dot product attention.

OVERRIDEABLE: The overridable backend for extension.

See `torch.nn.attention.sdpa_kernel()` for more details.

This class is in beta and subject to change.

