

GRU#

<https://pytorch.org/docs/stable/generated/torch.nn.GRU.html>

Description:

Apply a multi-layer gated recurrent unit (GRU) RNN to an input sequence. For each element in the input sequence, each layer computes the following function: where $h_{t,t}$ is the hidden state at time t , $x_{t,t}$ is the input at time t , $h_{(t-1)}$ is the hidden state of the layer at time $t-1$ or the initial hidden state at time 0, and $r_{t,t}$, $z_{t,t}$, $n_{t,t}$ are the reset, update, and new gates, respectively. σ is the sigmoid function, and $\odot$ is the Hadamard product. In a multilayer GRU, the input $x^{(l)}_{t,t(l)}$ of the l -th layer ($l \geq 2$) is the hidden state $h^{(l-1)}_{t,t(l-1)}$ of the previous layer multiplied by dropout $\delta^{(l-1)}_{t,t(l-1)}$ where each $\delta^{(l-1)}_{t,t(l-1)}$ is a Bernoulli random variable which is 0 with probability dropout.

input_size – The number of expected features in the input x
hidden_size – The number of features in the hidden state h

Function Signature

```
class torch.nn.GRU(input_size, hidden_size, num_layers=1,
bias=True, batch_first=False, dropout=0.0,
bidirectional=False, device=None, dtype=None)[source]#
```

Parameters

Inputs: input, h₀

Outputs: output, h_n

Variables

Code Examples

```
>>> rnn = nn.GRU(10, 20, 2)
>>> input = torch.randn(5, 3, 10)
>>> h0 = torch.randn(2, 3, 20)
>>> output, hn = rnn(input, h0)
```

Notes & Warnings

NOTE

All the weights and biases are initialized from $U(-k, k)$ where $k = \frac{1}{\sqrt{\text{hidden_size}}}$

NOTE

Note For bidirectional GRUs, forward and backward are directions 0 and 1 respectively. Example of splitting the output layers when `batch_first=False`: `output.view(seq_len, batch, num_directions, hidden_size)`.

NOTE

Note `batch_first` argument is ignored for unbatched inputs.

NOTE

Note The calculation of new gate `ntn_tnt` subtly differs from the original paper and other frameworks. In the original implementation, the Hadamard product $(\odot)$ between `rtr_trt` and the previous hidden state $h_{(t-1)}$ is done before the multiplication with the weight matrix W and addition of bias:

$$n_t = \tanh(W_{in} x_t + b_{in} + W_{hn} (r_t \odot h_{(t-1)}) + b_{hn})$$

This is in contrast to PyTorch implementation, which is done after $W_{hn} h_{(t-1)}$:

$$n_t = \tanh(W_{in} x_t + b_{in} + r_t \odot (W_{hn} h_{(t-1)} + b_{hn}))$$
NOTE

Note If the following conditions are satisfied: 1) cudnn is enabled, 2) input data is on the GPU 3) input data has dtype `torch.float16` 4) V100 GPU is used, 5) input data is not in `PackedSequence` format persistent algorithm can be selected to improve performance.

Detailed Documentation

Documentation

Apply a multi-layer gated recurrent unit (GRU) RNN to an input sequence. For each element in the input sequence, each layer computes the following function:

where $h_{t,t}$ is the hidden state at time t , $x_{t,t}$ is the input at time t , $h_{(t-1)}$ is the hidden state of the layer at time $t-1$ or the initial hidden state at time 0, and $r_{t,t}$, $z_{t,t}$, $n_{t,t}$ are the reset, update, and new gates, respectively. σ is the sigmoid function, and $\odot$ is the Hadamard product.

In a multilayer GRU, the input $x^{(l)}_{t,t}$ of the l -th layer ($l \geq 2$) is the hidden state $h^{(l-1)}_{t,t}$ of the previous layer multiplied by dropout $\delta^{(l-1)}_{t,t}$ where each $\delta^{(l-1)}_{t,t}$ is a Bernoulli random variable which is 0 or 1 with probability dropout.

input_size – The number of expected features in the input x

hidden_size – The number of features in the hidden state h

num_layers – Number of recurrent layers. E.g., setting `num_layers=2` would mean stacking two GRUs together to form a stacked GRU, with the second GRU taking in outputs of the first GRU and computing the final results. Default: 1

bias – If `False`, then the layer does not use bias weights b_{ih} and b_{hh} . Default: `True`

batch_first – If `True`, then the input and output tensors are provided as (batch, seq, feature) instead of (seq, batch, feature). Note that this does not apply to hidden or cell states. See the Inputs/Outputs sections below for details. Default: `False`

dropout – If non-zero, introduces a Dropout layer on the outputs of each GRU layer except the last layer, with dropout probability equal to dropout. Default: 0

bidirectional – If True, becomes a bidirectional GRU. Default: False

input: tensor of shape (L, H_{in}) for unbatched input, (L, N, H_{in}) when `batch_first=False` or (N, L, H_{in}) when `batch_first=True` containing the features of the input sequence. The input can also be a packed variable length sequence. See `torch.nn.utils.rnn.pack_padded_sequence()` or `torch.nn.utils.rnn.pack_sequence()` for details.

h_0 : tensor of shape $(D * \text{num_layers}, H_{out})$ or $(D * \text{num_layers}, N, H_{out})$ containing the initial hidden state for the input sequence. Defaults to zeros if not provided.

output: tensor of shape $(L, D * H_{out})$ for unbatched input, $(L, N, D * H_{out})$ when `batch_first=False` or $(N, L, D * H_{out})$ when `batch_first=True` containing the output features (h_t) from the last layer of the GRU, for each t . If a `torch.nn.utils.rnn.PackedSequence` has been given as the input, the output will also be a packed sequence.

h_n : tensor of shape $(D * \text{num_layers}, H_{out})$ or $(D * \text{num_layers}, N, H_{out})$ containing the final hidden state for the input sequence.

$\text{weight_ih_l}[k]$ – the learnable input-hidden weights of the k^{th} layer ($W_{ir}|W_{iz}|W_{in}$), of shape $(3 * \text{hidden_size}, \text{input_size})$ for $k = 0$. Otherwise, the shape is $(3 * \text{hidden_size}, \text{num_directions} * \text{hidden_size})$

$\text{weight_hh_l}[k]$ – the learnable hidden-hidden weights of the k^{th} layer ($W_{hr}|W_{hz}|W_{hn}$), of shape $(3 * \text{hidden_size}, \text{hidden_size})$

`bias_ih_l[k]` – the learnable input-hidden bias of the k^{th} layer (`b_ir|b_iz|b_in`), of shape $(3*\text{hidden_size})$

`bias_hh_l[k]` – the learnable hidden-hidden bias of the k^{th} layer (`b_hr|b_hz|b_hn`), of shape $(3*\text{hidden_size})$

All the weights and biases are initialized from $U(-k, k) \cdot \frac{1}{\sqrt{k}}$, $\frac{1}{\sqrt{k}}$ where $k = \frac{1}{\text{hidden_size}}$ for $k = 1$ to hidden_size

For bidirectional GRUs, forward and backward are directions 0 and 1 respectively. Example of splitting the output layers when `batch_first=False`: `output.view(seq_len, batch, num_directions, hidden_size)`.

`batch_first` argument is ignored for unbatched inputs.

The calculation of new gate `ntn_tnt` subtly differs from the original paper and other frameworks. In the original implementation, the Hadamard product $(\odot)$ between `rtr_trt` and the previous hidden state $h_{(t-1)}$ is done before the multiplication with the weight matrix W and addition of bias:

This is in contrast to PyTorch implementation, which is done after $W h_{(t-1)}$

This implementation differs on purpose for efficiency.

If the following conditions are satisfied: 1) cudnn is enabled, 2) input data is on the GPU 3) input data has dtype `torch.float16` 4) V100 GPU is used, 5) input data is not in `PackedSequence` format persistent algorithm can be selected to improve performance.