

Upsample#

<https://pytorch.org/docs/stable/generated/torch.nn.Upsample.html>

Description:

Upsamples a given multi-channel 1D (temporal), 2D (spatial) or 3D (volumetric) data. The input data is assumed to be of the form minibatch x channels x [optional depth] x [optional height] x width. Hence, for spatial inputs, we expect a 4D Tensor and for volumetric inputs, we expect a 5D Tensor. The algorithms available for upsampling are nearest neighbor and linear, bilinear, bicubic and trilinear for 3D, 4D and 5D input Tensor, respectively. One can either give a `scale_factor` or the target output size to calculate the output size. (You cannot give both, as it is ambiguous) size (int or Tuple[int] or Tuple[int, int] or Tuple[int, int, int], optional) – output spatial sizes

Function Signature

```
class torch.nn.Upsample(size=None, scale_factor=None,
mode='nearest', align_corners=None,
recompute_scale_factor=None)[source]#
```

Parameters

Shape:

`extra_repr()[source]`#

Return type

Returns:

Tensor

Code Examples

```
>>> input = torch.arange(1, 5, dtype=torch.float32).view(1, 1,
2, 2)
>>> input
tensor([[[[1., 2.],
          [3., 4.]]]])

>>> m = nn.Upsample(scale_factor=2, mode='nearest')
>>> m(input)
tensor([[[[1., 1., 2., 2.],
          [1., 1., 2., 2.]],
```

```
[3., 3., 4., 4.],  
[3., 3., 4., 4.]])])])
```

```
>>> m = nn.Upsample(scale_factor=2, mode='bilinear') #  
align_corners=False
```

```
>>> m(input)  
tensor([[[[1.0000, 1.2500, 1.7500, 2.0000],  
          [1.5000, 1.7500, 2.2500, 2.5000],  
          [2.5000, 2.7500, 3.2500, 3.5000],  
          [3.0000, 3.2500, 3.7500, 4.0000]]]])])
```

```
>>> m = nn.Upsample(scale_factor=2, mode='bilinear',  
align_corners=True)
```

```
>>> m(input)  
tensor([[[[1.0000, 1.3333, 1.6667, 2.0000],  
          [1.6667, 2.0000, 2.3333, 2.6667],  
          [2.3333, 2.6667, 3.0000, 3.3333],  
          [3.0000, 3.3333, 3.6667, 4.0000]]]])])
```

```
>>> # Try scaling the same data in a larger tensor
```

```
>>> input_3x3 = torch.zeros(3, 3).view(1, 1, 3, 3)
```

```
>>> input_3x3[:, :, :2, :2].copy_(input)
```

```
tensor([[[[1., 2.,  
          [3., 4.]])]])])
```

```
>>> input_3x3
```

```
tensor([[[[1., 2., 0.,  
          [3., 4., 0.,  
          [0., 0., 0.]])]])])
```

```
>>> m = nn.Upsample(scale_factor=2, mode='bilinear') #  
align_corners=False
```

```
>>> # Notice that values in top left corner are the same with  
the small input (except at boundary)
```

```
>>> m(input_3x3)  
tensor([[[[1.0000, 1.2500, 1.7500, 1.5000, 0.5000, 0.0000],  
          [1.5000, 1.7500, 2.2500, 1.8750, 0.6250, 0.0000],  
          [2.5000, 2.7500, 3.2500, 2.6250, 0.8750, 0.0000],  
          [2.2500, 2.4375, 2.8125, 2.2500, 0.7500, 0.0000],
```

```

        [0.7500, 0.8125, 0.9375, 0.7500, 0.2500, 0.0000],
        [0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000]]]])

>>> m = nn.Upsample(scale_factor=2, mode='bilinear',
align_corners=True)
>>> # Notice that values in top left corner are now changed
>>> m(input_3x3)
tensor([[[[1.0000, 1.4000, 1.8000, 1.6000, 0.8000, 0.0000],
          [1.8000, 2.2000, 2.6000, 2.2400, 1.1200, 0.0000],
          [2.6000, 3.0000, 3.4000, 2.8800, 1.4400, 0.0000],
          [2.4000, 2.7200, 3.0400, 2.5600, 1.2800, 0.0000],
          [1.2000, 1.3600, 1.5200, 1.2800, 0.6400, 0.0000],
          [0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000]]]])

```

Notes & Warnings

⚠ WARNING

Warning With `align_corners = True`, the linearly interpolating modes (linear, bilinear, bicubic, and trilinear) don't proportionally align the output and input pixels, and thus the output values can depend on the input size. This was the default behavior for these modes up to version 0.3.1. Since then, the default behavior is `align_corners = False`. See below for concrete examples on how this affects the outputs.

NOTE

Note If you want downsampling/general resizing, you should use `interpolate()`.

Detailed Documentation

Documentation

Upsamples a given multi-channel 1D (temporal), 2D (spatial) or 3D (volumetric) data.

The input data is assumed to be of the form minibatch x channels x [optional depth] x [optional height] x width. Hence, for spatial inputs, we expect a 4D Tensor and for volumetric inputs, we expect a 5D Tensor.

The algorithms available for upsampling are nearest neighbor and linear, bilinear, bicubic and trilinear for 3D, 4D and 5D input Tensor, respectively.

One can either give a `scale_factor` or the target output size to calculate the output size. (You cannot give both, as it is ambiguous)

`size` (int or Tuple[int] or Tuple[int, int] or Tuple[int, int, int], optional) – output spatial sizes

`scale_factor` (float or Tuple[float] or Tuple[float, float] or Tuple[float, float, float], optional) – multiplier for spatial size. Has to match input size if it is a tuple.

`mode` (str, optional) – the upsampling algorithm: one of 'nearest', 'linear', 'bilinear', 'bicubic' and 'trilinear'. Default: 'nearest'

`align_corners` (bool, optional) – if True, the corner pixels of the input and output tensors are aligned, and thus preserving the values at those pixels. This only has effect when mode is 'linear', 'bilinear', 'bicubic', or 'trilinear'. Default: False

`recompute_scale_factor` (bool, optional) – recompute the `scale_factor` for use in the interpolation calculation. If `recompute_scale_factor` is True, then `scale_factor` must be passed in and `scale_factor` is used to compute the output size. The computed output size will be used to infer new scales for the interpolation. Note that when `scale_factor` is

floating-point, it may differ from the recomputed `scale_factor` due to rounding and precision issues. If `recompute_scale_factor` is `False`, then `size` or `scale_factor` will be used directly for interpolation.

Input: (N, C, W_{in}) (N, C, W_{in}) , (N, C, H_{in}, W_{in}) (N, C, H_{in}, W_{in}) or $(N, C, D_{in}, H_{in}, W_{in})$ $(N, C, D_{in}, H_{in}, W_{in})$

Output: (N, C, W_{out}) (N, C, W_{out}) , (N, C, H_{out}, W_{out}) (N, C, H_{out}, W_{out}) or $(N, C, D_{out}, H_{out}, W_{out})$ $(N, C, D_{out}, H_{out}, W_{out})$, where

With `align_corners = True`, the linearly interpolating modes (linear, bilinear, bicubic, and trilinear) don't proportionally align the output and input pixels, and thus the output values can depend on the input size. This was the default behavior for these modes up to version 0.3.1. Since then, the default behavior is `align_corners = False`. See below for concrete examples on how this affects the outputs.

If you want downsampling/general resizing, you should use `interpolate()`.

Return the extra representation of the module.

Runs the forward pass.