

LBFGS#

<https://pytorch.org/docs/stable/generated/torch.optim.lbfgs.LBFGS.html>

Description:

Implements L-BFGS algorithm. Heavily inspired by minFunc. This optimizer doesn't support per-parameter options and parameter groups (there can be only one).

Function Signature

```
class torch.optim.lbfgs.LBFGS(params, lr=1, max_iter=20,
max_eval=None, tolerance_grad=1e-07, tolerance_change=1e-09,
history_size=100, line_search_fn=None)[source]#
```

Parameters

```
add_param_group(param_group)[source]#
```

Parameters

```
load_state_dict(state_dict)[source]#
```

Returns:

dict[str, Any]

Code Examples

```
>>> model = torch.nn.Linear(10, 10)
>>> optim = torch.optim.SGD(model.parameters(), lr=3e-4)
>>> scheduler1 = torch.optim.lr_scheduler.LinearLR(
...     optim,
...     start_factor=0.1,
...     end_factor=1,
...     total_iters=20,
... )
>>> scheduler2 = torch.optim.lr_scheduler.CosineAnnealingLR(
...     optim,
...     T_max=80,
...     eta_min=3e-5,
... )
>>> lr = torch.optim.lr_scheduler.SequentialLR(
...     optim,
...     schedulers=[scheduler1, scheduler2],
...     milestones=[20],
... )
>>> lr.load_state_dict(torch.load("./save_seq.pt"))
>>> # now load the optimizer checkpoint after loading the
LRScheduler
>>> optim.load_state_dict(torch.load("./save_optim.pt"))
```

```
hook(optimizer) -> None
```

```
hook(optimizer, state_dict) -> state_dict or None
```

```
hook(optimizer, state_dict) -> state_dict or None
```

```
hook(optimizer) -> None
```

```
hook(optimizer, args, kwargs) -> None
```

```
hook(optimizer, args, kwargs) -> None or modified args and  
kwargs
```

```

{
  'state': {
    0: {'momentum_buffer': tensor(...), ...},
    1: {'momentum_buffer': tensor(...), ...},
    2: {'momentum_buffer': tensor(...), ...},
    3: {'momentum_buffer': tensor(...), ...}
  },
  'param_groups': [
    {
      'lr': 0.01,
      'weight_decay': 0,
      ...
      'params': [0]
      'param_names' ['param0'] (optional)
    },
    {
      'lr': 0.001,
      'weight_decay': 0.5,
      ...
      'params': [1, 2, 3]
      'param_names': ['param1', 'layer.weight',
'layer.bias'] (optional)
    }
  ]
}

```

Notes & Warnings

 **WARNING** This optimizer doesn't support per-parameter options and parameter groups (there can be only one).

WARNING

Warning Right now all parameters have to be on a single device. This will be improved in the future.

NOTE

Note This is a very memory intensive optimizer (it requires additional $\text{param_bytes} * (\text{history_size} + 1)$ bytes). If it doesn't fit in memory try reducing the history size, or use a different algorithm.

WARNING

Warning Make sure this method is called after initializing `torch.optim.lr_scheduler.LRScheduler`, as calling it beforehand will overwrite the loaded learning rates.

NOTE

Note The names of the parameters (if they exist under the “`param_names`” key of each param group in `state_dict()`) will not affect the loading process. To use the parameters' names for custom cases (such as when the parameters in the loaded state dict differ from those initialized in the optimizer), a custom `register_load_state_dict_pre_hook` should be implemented to adapt the loaded dict accordingly. If `param_names` exist in loaded state dict `param_groups` they will be saved and override the current names, if present, in the optimizer state. If they do not exist in loaded state dict, the optimizer `param_names` will remain unchanged.

Detailed Documentation

Documentation

Implements L-BFGS algorithm.

Heavily inspired by minFunc.

This optimizer doesn't support per-parameter options and parameter groups (there can be only one).

Right now all parameters have to be on a single device. This will be improved in the future.

This is a very memory intensive optimizer (it requires additional $\text{param_bytes} * (\text{history_size} + 1)$ bytes). If it doesn't fit in memory try reducing the history size, or use a different algorithm.

`params` (iterable) – iterable of parameters to optimize. Parameters must be real.

`lr` (float, optional) – learning rate (default: 1)

`max_iter` (int, optional) – maximal number of iterations per optimization step (default: 20)

`max_eval` (int, optional) – maximal number of function evaluations per optimization step (default: $\text{max_iter} * 1.25$).

`tolerance_grad` (float, optional) – termination tolerance on first order optimality (default: $1e-7$).

`tolerance_change` (float, optional) – termination tolerance on function value/parameter changes (default: $1e-9$).

`history_size` (int, optional) – update history size (default: 100).

`line_search_fn` (str, optional) – either ‘strong_wolfe’ or None (default: None).

Add a param group to the Optimizer’s `param_groups`.

This can be useful when fine tuning a pre-trained network as frozen layers can be made trainable and added to the Optimizer as training progresses.

`param_group` (dict) – Specifies what Tensors should be optimized along with group specific optimization options.

Load the optimizer state.

`state_dict` (dict) – optimizer state. Should be an object returned from a call to `state_dict()`.

Make sure this method is called after initializing `torch.optim.lr_scheduler.LRScheduler`, as calling it beforehand will overwrite the loaded learning rates.

The names of the parameters (if they exist under the “`param_names`” key of each param group in `state_dict()`) will not affect the loading process. To use the parameters’ names for custom cases (such as when the parameters in the loaded state dict differ from those initialized in the optimizer), a custom `register_load_state_dict_pre_hook` should be implemented to adapt the loaded dict accordingly. If `param_names` exist in loaded state dict `param_groups` they will be saved and override the current names, if present, in the optimizer state. If they do not exist in loaded state dict, the optimizer `param_names` will remain unchanged.

Register a `load_state_dict` post-hook which will be called after `load_state_dict()` is called. It should have the following signature:

The optimizer argument is the optimizer instance being used.

The hook will be called with argument `self` after calling `load_state_dict` on `self`. The registered hook can be used to perform post-processing after `load_state_dict` has loaded the `state_dict`.

`hook (Callable)` – The user defined hook to be registered.

`prepend (bool)` – If `True`, the provided post hook will be fired before all the already registered post-hooks on `load_state_dict`. Otherwise, the provided hook will be fired after all the already registered post-hooks. (default: `False`)

a handle that can be used to remove the added hook by calling `handle.remove()`

`torch.utils.hooks.RemoveableHandle`

Register a `load_state_dict` pre-hook which will be called before `load_state_dict()` is called. It should have the following signature:

The `optimizer` argument is the optimizer instance being used and the `state_dict` argument is a shallow copy of the `state_dict` the user passed in to `load_state_dict`. The hook may modify the `state_dict` inplace or optionally return a new one. If a `state_dict` is returned, it will be used to be loaded into the optimizer.

The hook will be called with argument `self` and `state_dict` before calling `load_state_dict` on `self`. The registered hook can be used to perform pre-processing before the `load_state_dict` call is made.

`hook (Callable)` – The user defined hook to be registered.

`prepend (bool)` – If `True`, the provided pre hook will be fired before all the already registered pre-hooks on `load_state_dict`. Otherwise, the provided hook will be fired after all the already registered pre-hooks. (default: `False`)

a handle that can be used to remove the added hook by calling `handle.remove()`

`torch.utils.hooks.RemoveableHandle`

Register a state dict post-hook which will be called after `state_dict()` is called.

It should have the following signature:

The hook will be called with arguments `self` and `state_dict` after generating a `state_dict` on `self`. The hook may modify the `state_dict` inplace or optionally return a new one. The registered hook can be used to perform post-processing on the `state_dict` before it is returned.

`hook (Callable)` – The user defined hook to be registered.

`prepend (bool)` – If `True`, the provided post hook will be fired before all the already registered post-hooks on `state_dict`. Otherwise, the provided hook will be fired after all the already registered post-hooks. (default: `False`)

a handle that can be used to remove the added hook by calling `handle.remove()`

`torch.utils.hooks.RemoveableHandle`

Register a state dict pre-hook which will be called before `state_dict()` is called.

It should have the following signature:

The optimizer argument is the optimizer instance being used. The hook will be called with argument `self` before calling `state_dict` on `self`. The registered hook can be used to perform pre-processing before the `state_dict` call is made.

`hook (Callable)` – The user defined hook to be registered.

prepend (bool) – If True, the provided pre hook will be fired before all the already registered pre-hooks on state_dict. Otherwise, the provided hook will be fired after all the already registered pre-hooks. (default: False)

a handle that can be used to remove the added hook by calling handle.remove()

torch.utils.hooks.RemoveableHandle

Register an optimizer step post hook which will be called after optimizer step.

It should have the following signature:

The optimizer argument is the optimizer instance being used.

hook (Callable) – The user defined hook to be registered.

a handle that can be used to remove the added hook by calling handle.remove()

torch.utils.hooks.RemovableHandle

Register an optimizer step pre hook which will be called before optimizer step.

It should have the following signature:

The optimizer argument is the optimizer instance being used. If args and kwargs are modified by the pre-hook, then the transformed values are returned as a tuple containing the new_args and new_kwargs.

hook (Callable) – The user defined hook to be registered.

a handle that can be used to remove the added hook by calling handle.remove()

torch.utils.hooks.RemovableHandle

Return the state of the optimizer as a dict.

It contains two entries:

differs between optimizer classes, but some common characteristics hold. For example, state is saved per parameter, and the parameter itself is NOT saved. state is a Dictionary mapping parameter ids to a Dict with state corresponding to each parameter.

parameter group is a Dict. Each parameter group contains metadata specific to the optimizer, such as learning rate and weight decay, as well as a List of parameter IDs of the parameters in the group. If a param group was initialized with `named_parameters()` the names content will also be saved in the state dict.

NOTE: The parameter IDs may look like indices but they are just IDs associating state with `param_group`. When loading from a `state_dict`, the optimizer will zip the `param_group` params (int IDs) and the optimizer `param_groups` (actual `nn.Parameter s`) in order to match state WITHOUT additional verification.

A returned state dict might look something like:

Perform a single optimization step.

`closure` (Callable) – A closure that reevaluates the model and returns the loss.

Reset the gradients of all optimized `torch.Tensor s`.

`set_to_none` (bool, optional) –

Instead of setting to zero, set the grads to None. Default: True

This will in general have lower memory footprint, and can modestly improve performance. However, it changes certain behaviors. For example:

When the user tries to access a gradient and perform manual ops on it, a `None` attribute or a Tensor full of 0s will behave differently.

If the user requests `zero_grad(set_to_none=True)` followed by a backward pass, `.grads` are guaranteed to be `None` for params that did not receive a gradient.

`torch.optim` optimizers have a different behavior if the gradient is 0 or `None` (in one case it does the step with a gradient of 0 and in the other it skips the step altogether).