

StepLR#

[https://pytorch.org/docs/stable/generated/torch.optim.lr_scheduler.StepLR.ht](https://pytorch.org/docs/stable/generated/torch.optim.lr_scheduler.StepLR.html)

Description:

Decays the learning rate of each parameter group by gamma every step_size epochs. Notice that such decay can happen simultaneously with other changes to the learning rate from outside this scheduler. When last_epoch=-1, sets initial lr as lr. optimizer (Optimizer) – Wrapped optimizer. step_size (int) – Period of learning rate decay. gamma (float) – Multiplicative factor of learning rate decay. Default: 0.1.

Function Signature

```
class torch.optim.lr_scheduler.StepLR(optimizer, step_size, gamma=0.1, last_epoch=-1)[source]#
```

Parameters

```
get_last_lr()[source]#
```

Return type

```
get_lr()[source]#
```

Returns:

dict[str, Any]

Code Examples

```
>>> # Assuming optimizer uses lr = 0.05 for all groups
>>> # lr = 0.05      if epoch < 30
>>> # lr = 0.005     if 30 <= epoch < 60
>>> # lr = 0.0005    if 60 <= epoch < 90
>>> # ...
>>> scheduler = StepLR(optimizer, step_size=30, gamma=0.1)
>>> for epoch in range(100):
>>>     train(...)
>>>     validate(...)
>>>     scheduler.step()
```

Detailed Documentation

Documentation

Decays the learning rate of each parameter group by gamma every step_size epochs.

Notice that such decay can happen simultaneously with other changes to the learning rate from outside this scheduler. When last_epoch=-1, sets initial lr as lr.

optimizer (Optimizer) – Wrapped optimizer.

step_size (int) – Period of learning rate decay.

gamma (float) – Multiplicative factor of learning rate decay. Default: 0.1.

last_epoch (int) – The index of last epoch. Default: -1.

Return last computed learning rate by current scheduler.

Compute the learning rate of each parameter group.

Load the scheduler's state.

state_dict (dict) – scheduler state. Should be an object returned from a call to state_dict().

Return the state of the scheduler as a dict.

It contains an entry for every variable in self.__dict__ which is not the optimizer.
