

PerAxis#

<https://pytorch.org/docs/stable/generated/torch.ao.quantization.observer.Per>

Description:

Represents per-axis granularity in quantization. This granularity type calculates different quantization parameters along a specified axis of the tensor. For example if the input tensor is shape [8, 16] and axis=0, then the quantization parameters are calculated for each row of the tensor. Giving a total of 8 quantization parameters. axis (int) – The axis along which reduction is performed.

Function Signature

```
class torch.ao.quantization.observer.PerAxis(axis)[source]#  
Variables
```

Detailed Documentation

Documentation

Represents per-axis granularity in quantization.

This granularity type calculates different quantization parameters along a specified axis of the tensor.

For example if the input tensor is shape [8, 16] and axis=0, then the quantization parameters are calculated for each row of the tensor. Giving a total of 8 quantization parameters.

axis (int) – The axis along which reduction is performed.

Represents per-axis granularity in quantization.

This granularity type calculates different quantization parameters along a specified axis of the tensor.

For example if the input tensor is shape [8, 16] and axis=0, then the quantization parameters are calculated for each row of the tensor. Giving a total of 8 quantization parameters.

axis (int) – The axis along which reduction is performed.