

PerRow#

<https://pytorch.org/docs/stable/generated/torch.ao.quantization.observer.Per>

Description:

Represents row-wise granularity in quantization. This is a special case of per-axis quantization and is unique to Float8 matmuls where the input is quantized with a block_size of (1, ..., input.shape[-1]). And the weight is quantized with a block_size of (1, weight.shape[1]).

Function Signature

```
class torch.ao.quantization.observer.PerRow[source]#
```

Detailed Documentation

Documentation

Represents row-wise granularity in quantization.

This is a special case of per-axis quantization and is unique to Float8 matmuls where the input is quantized with a `block_size` of `(1, ..., input.shape[-1])`. And the weight is quantized with a `block_size` of `(1, weight.shape[1])`.

Represents row-wise granularity in quantization.

This is a special case of per-axis quantization and is unique to Float8 matmuls where the input is quantized with a `block_size` of `(1, ..., input.shape[-1])`. And the weight is quantized with a `block_size` of `(1, weight.shape[1])`.