

QConfig#

<https://pytorch.org/docs/stable/generated/torch.ao.quantization.qconfig.QConfig>

Description:

Describes how to quantize a layer or a part of the network by providing settings (observer classes) for activations and weights respectively. Note that QConfig needs to contain observer classes (like MinMaxObserver) or a callable that returns instances on invocation, not the concrete observer instances themselves. Quantization preparation function will instantiate observers multiple times for each of the layers. Observer classes have usually reasonable default arguments, but they can be overwritten with `with_args` method (that behaves like `functools.partial`):

Function Signature

```
class torch.ao.quantization.qconfig.QConfig(activation,  
weight)[source]#
```

Code Examples

```
my_qconfig = QConfig(  
    activation=MinMaxObserver.with_args(dtype=torch.qint8),  
    weight=default_observer.with_args(dtype=torch.qint8),  
)
```

Detailed Documentation

Documentation

Describes how to quantize a layer or a part of the network by providing settings (observer classes) for activations and weights respectively.

Note that `QConfig` needs to contain observer classes (like `MinMaxObserver`) or a callable that returns instances on invocation, not the concrete observer instances themselves. Quantization preparation function will instantiate observers multiple times for each of the layers.

Observer classes have usually reasonable default arguments, but they can be overwritten with `with_args` method (that behaves like `functools.partial`):