

quantize_dynamic#

https://pytorch.org/docs/stable/generated/torch.ao.quantization.quantize_dynamic.html

Description:

Converts a float model to dynamic (i.e. weights-only) quantized model. Replaces specified modules with dynamic weight-only quantized versions and output the quantized model. For simplest usage provide dtype argument that can be float16 or qint8. Weight-only quantization by default is performed for layers with large weights size - i.e. Linear and RNN variants. Fine grained control is possible with qconfig and mapping that act similarly to quantize(). If qconfig is provided, the dtype argument is ignored.

Function Signature

```
class torch.ao.quantization.quantize_dynamic(model,  
qconfig_spec=None, dtype=torch.qint8, mapping=None,  
inplace=False)[source]#
```

Parameters

Detailed Documentation

Documentation

Converts a float model to dynamic (i.e. weights-only) quantized model.

Replaces specified modules with dynamic weight-only quantized versions and output the quantized model.

For simplest usage provide dtype argument that can be float16 or qint8. Weight-only quantization by default is performed for layers with large weights size - i.e. Linear and RNN variants.

Fine grained control is possible with qconfig and mapping that act similarly to quantize(). If qconfig is provided, the dtype argument is ignored.

A dictionary that maps from name or type of submodule to quantization configuration, qconfig applies to all submodules of a given module unless qconfig for the submodules are specified (when the submodule already has qconfig attribute). Entries in the dictionary need to be QConfig instances.

A set of types and/or submodule names to apply dynamic quantization to, in which case the dtype argument is used to specify the bit-width

inplace – carry out model transformations in-place, the original module is mutated

mapping – maps type of a submodule to a type of corresponding dynamically quantized version with which the submodule needs to be replaced
