

torch.nanquantile#

<https://pytorch.org/docs/stable/generated/torch.nanquantile.html#torch.nanquantile>

Description:

This is a variant of `torch.quantile()` that “ignores” NaN values, computing the quantiles `q` as if NaN values in input did not exist. If all values in a reduced row are NaN then the quantiles for that reduction will be NaN. See the documentation for `torch.quantile()`. `input` (Tensor) – the input tensor. `q` (float or Tensor) – a scalar or 1D tensor of quantile values in the range `[0, 1]` `dim` (int, optional) – the dimension to reduce. If `None`, all dimensions are reduced. `keepdim` (bool, optional) – whether the output tensor has dim retained or not. Default: `False`.

Function Signature

```
torch.nanquantile(input, q, dim=None, keepdim=False, *,  
interpolation='linear', out=None) → Tensor#
```

Parameters

Keyword Arguments

Parameters

Keyword

interpolation (str) – interpolation method to use when the desired quantile lies between two data points. Can be linear, lower, higher, midpoint and nearest. Default is linear. out (Tensor, optional) – the output tensor.

Code Examples

```
>>> t = torch.tensor([float('nan'), 1, 2])
>>> t.quantile(0.5)
tensor(nan)
>>> t.nanquantile(0.5)
tensor(1.5000)
>>> t = torch.tensor([[float('nan'), float('nan')], [1, 2]])
>>> t
tensor([[nan, nan],
        [1., 2.]])
>>> t.nanquantile(0.5, dim=0)
tensor([1., 2.])
>>> t.nanquantile(0.5, dim=1)
tensor([ nan, 1.5000])
```

Detailed Documentation

Documentation

This is a variant of `torch.quantile()` that “ignores” NaN values, computing the quantiles `q` as if NaN values in input did not exist. If all values in a reduced row are NaN then the quantiles for that reduction will be NaN. See the documentation for `torch.quantile()`.

`input (Tensor)` – the input tensor.

`q (float or Tensor)` – a scalar or 1D tensor of quantile values in the range `[0, 1]`

`dim (int, optional)` – the dimension to reduce. If `None`, all dimensions are reduced.

`keepdim (bool, optional)` – whether the output tensor has `dim` retained or not. Default: `False`.

`interpolation (str)` – interpolation method to use when the desired quantile lies between two data points. Can be `linear`, `lower`, `higher`, `midpoint` and `nearest`. Default is `linear`.

`out (Tensor, optional)` – the output tensor.