

SVC#

<https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html>

Description:

C-Support Vector Classification. The implementation is based on libsvm. The fit time scales at least quadratically with the number of samples and may be impractical beyond tens of thousands of samples. For large datasets consider using LinearSVC or SGDClassifier instead, possibly after a Nystroem transformer or other Kernel Approximation. The multiclass support is handled according to a one-vs-one scheme. For details on the precise mathematical formulation of the provided kernel functions and how gamma, coef0 and degree affect each other, see the corresponding section in the narrative documentation: Kernel functions.

Function Signature

```
class sklearn.svm.SVC(*, C=1.0, kernel='rbf', degree=3,
gamma='scale', coef0=0.0, shrinking=True, probability=False,
tol=0.001, cache_size=200, class_weight=None, verbose=False,
max_iter=-1, decision_function_shape='ovr',
break_ties=False, random_state=None)[source]#
```

Parameters:

Cfloat, default=1.0

kernel{'linear', 'poly', 'rbf', 'sigmoid', 'precomputed'} or callable, default='rbf'

degreeint, default=3

```
class sklearn.svm.SVC(*, C=1.0, kernel='rbf', degree=3,
gamma='scale', coef0=0.0, shrinking=True, probability=False,
tol=0.001, cache_size=200, class_weight=None, verbose=False,
max_iter=-1, decision_function_shape='ovr',
break_ties=False, random_state=None)[source]#
```

Parameters:

Cfloat, default=1.0

kernel{'linear', 'poly', 'rbf', 'sigmoid', 'precomputed'} or callable, default='rbf'

degreeint, default=3

decision_function(X)[source]#

Parameters:

Xarray-like of shape (n_samples, n_features)

Returns:

Xndarray of shape (n_samples, n_classes * (n_classes-1) / 2)

Parameters

tolfloat, default=1e

Tolerance for stopping criterion.

`max_iterint, default=`

Hard limit on iterations within solver, or -1 for no limit.

`coef_ndarray of shape (n_classes * (n_classes`

Weights assigned to the features when kernel="linear".

`dual_coef_ndarray of shape (n_classes`

Dual coefficients of the support vector in the decision function (see Mathematical formulation), multiplied by their targets. For multiclass, coefficient for all 1-vs-1 classifiers. The layout of the coefficients in the multiclass case is somewhat non-trivial. See the multi-class section of the User Guide for details.

`intercept_ndarray of shape (n_classes * (n_classes`

Constants in decision function.

`n_iter_ndarray of shape (n_classes * (n_classes`

Number of iterations run by the optimization routine to fit the model. The shape of this attribute depends on the number of models optimized which in turn depends on the number of classes. Added in version 1.1.

probA_ndarray of shape (n_classes * (n_classes

Parameter learned in Platt scaling when probability=True.

probB_ndarray of shape (n_classes * (n_classes

Parameter learned in Platt scaling when probability=True.

Xarray

The input samples.

Xndarray of shape (n_samples, n_classes * (n_classes

Returns the decision function of the sample for each class in the model. If decision_function_shape='ovr', the shape is (n_samples, n_classes).

X{array

Training vectors, where n_samples is the number of samples and n_features is the number of features. For kernel="precomputed", the expected shape of X is (n_samples, n_samples).

`yarray`

Target values (class labels in classification, real numbers in regression).

`sample_weightarray`

Per-sample weights. Rescale C per sample. Higher weights force the classifier to put more emphasis on these points.

`paramsdict`

Parameter names mapped to their values.

`X{array`

For `kernel="precomputed"`, the expected shape of X is (n_samples_test, n_samples_train).

`Xarray`

For `kernel="precomputed"`, the expected shape of X is (n_samples_test, n_samples_train).

`Xarray`

For `kernel="precomputed"`, the expected shape of `X` is `(n_samples_test, n_samples_train)`.

`Xarray`

Test samples.

`yarray`

True labels for `X`.

`sample_weightarray`

Sample weights.

`**paramsdict`

Estimator parameters.

Code Examples

```
>>> import numpy as np
>>> from sklearn.pipeline import make_pipeline
>>> from sklearn.preprocessing import StandardScaler
>>> X = np.array([[-1, -1], [-2, -1], [1, 1], [2, 1]])
>>> y = np.array([1, 1, 2, 2])
>>> from sklearn.svm import SVC
>>> clf = make_pipeline(StandardScaler(), SVC(gamma='auto'))
>>> clf.fit(X, y)
Pipeline(steps=[('standardscaler', StandardScaler()),
                 ('svc', SVC(gamma='auto'))])
```

```
>>> print(clf.predict([[-0.8, -1]]))
[1]
```

Notes & Warnings

NOTE See also SVRSupport Vector Machine for Regression implemented using libsvm. LinearSVCSscalable Linear Support Vector Machine for classification implemented using liblinear. Check the See Also section of LinearSVC for more comparison element.

Detailed Documentation

Documentation

C-Support Vector Classification.

The implementation is based on libsvm. The fit time scales at least quadratically with the number of samples and may be impractical beyond tens of thousands of samples. For large datasets consider using LinearSVC or SGDClassifier instead, possibly after a Nystroem transformer or other Kernel Approximation.

The multiclass support is handled according to a one-vs-one scheme.

For details on the precise mathematical formulation of the provided kernel functions and how gamma, coef0 and degree affect each other, see the corresponding section in the narrative documentation: Kernel functions.

To learn how to tune SVC's hyperparameters, see the following example: Nested versus non-nested cross-validation

Read more in the User Guide.

Regularization parameter. The strength of the regularization is inversely proportional to C. Must be strictly positive. The penalty is a squared l_2 penalty. For an intuitive visualization of the effects of scaling the regularization parameter C, see Scaling the regularization parameter for SVCs.

Specifies the kernel type to be used in the algorithm. If none is given, 'rbf' will be used. If a callable is given it is used to pre-compute the kernel matrix from data matrices; that matrix should be an array of shape (n_samples, n_samples). For an intuitive visualization of different kernel types see Plot classification boundaries with different SVM Kernels.

Degree of the polynomial kernel function ('poly'). Must be non-negative. Ignored by all

other kernels.

Kernel coefficient for 'rbf', 'poly' and 'sigmoid'.

if `gamma='scale'` (default) is passed then it uses $1 / (n_features * X.var())$ as value of `gamma`,

if 'auto', uses $1 / n_features$

if float, must be non-negative.

Changed in version 0.22: The default value of `gamma` changed from 'auto' to 'scale'.

Independent term in kernel function. It is only significant in 'poly' and 'sigmoid'.

Whether to use the shrinking heuristic. See the User Guide.

Whether to enable probability estimates. This must be enabled prior to calling `fit`, will slow down that method as it internally uses 5-fold cross-validation, and `predict_proba` may be inconsistent with `predict`. Read more in the User Guide.

Tolerance for stopping criterion.

Specify the size of the kernel cache (in MB).

Set the parameter `C` of class `i` to `class_weight[i]*C` for SVC. If not given, all classes are supposed to have weight one. The "balanced" mode uses the values of `y` to automatically adjust weights inversely proportional to class frequencies in the input data as $n_samples / (n_classes * np.bincount(y))$.

Enable verbose output. Note that this setting takes advantage of a per-process runtime setting in `libsvm` that, if enabled, may not work properly in a multithreaded context.

Hard limit on iterations within solver, or -1 for no limit.

Whether to return a one-vs-rest ('ovr') decision function of shape (n_samples, n_classes) as all other classifiers, or the original one-vs-one ('ovo') decision function of libsvm which has shape (n_samples, n_classes * (n_classes - 1) / 2). However, note that internally, one-vs-one ('ovo') is always used as a multi-class strategy to train models; an ovr matrix is only constructed from the ovo matrix. The parameter is ignored for binary classification.

Changed in version 0.19: decision_function_shape is 'ovr' by default.

Added in version 0.17: decision_function_shape='ovr' is recommended.

Changed in version 0.17: Deprecated decision_function_shape='ovo' and None.

If true, decision_function_shape='ovr', and number of classes > 2, predict will break ties according to the confidence values of decision_function; otherwise the first class among the tied classes is returned. Please note that breaking ties comes at a relatively high computational cost compared to a simple predict. See SVM Tie Breaking Example for an example of its usage with decision_function_shape='ovr'.

Added in version 0.22.

Controls the pseudo random number generation for shuffling the data for probability estimates. Ignored when probability is False. Pass an int for reproducible output across multiple function calls. See Glossary.

Multipliers of parameter C for each class. Computed based on the class_weight parameter.

Weights assigned to the features when kernel="linear".

Dual coefficients of the support vector in the decision function (see Mathematical formulation), multiplied by their targets. For multiclass, coefficient for all 1-vs-1 classifiers. The layout of the coefficients in the multiclass case is somewhat non-trivial. See the multi-class section of the User Guide for details.

0 if correctly fitted, 1 otherwise (will raise warning)

Constants in decision function.

Number of features seen during fit.

Added in version 0.24.

Names of features seen during fit. Defined only when X has feature names that are all strings.

Added in version 1.0.

Number of iterations run by the optimization routine to fit the model. The shape of this attribute depends on the number of models optimized which in turn depends on the number of classes.

Added in version 1.1.

Indices of support vectors.

Support vectors. An empty array if kernel is precomputed.

Number of support vectors for each class.

Parameter learned in Platt scaling when probability=True.

Parameter learned in Platt scaling when probability=True.

Array dimensions of training vector X.

Support Vector Machine for Regression implemented using libsvm.

Scalable Linear Support Vector Machine for classification implemented using liblinear.
Check the See Also section of LinearSVC for more comparison element.

LIBSVM: A Library for Support Vector Machines

Platt, John (1999). “Probabilistic Outputs for Support Vector Machines and Comparisons to Regularized Likelihood Methods”

For a comparison of the SVC with other classifiers see: Plot classification probability.

Evaluate the decision function for the samples in X.

Returns the decision function of the sample for each class in the model. If `decision_function_shape='ovr'`, the shape is (n_samples, n_classes).

If `decision_function_shape='ovo'`, the function values are proportional to the distance of the samples X to the separating hyperplane. If the exact distances are required, divide the function values by the norm of the weight vector (`coef_`). See also this question for further details. If `decision_function_shape='ovr'`, the decision function is a monotonic transformation of ovo decision function.

Fit the SVM model according to the given training data.

Training vectors, where n_samples is the number of samples and n_features is the number of features. For `kernel="precomputed"`, the expected shape of X is (n_samples, n_samples).

Target values (class labels in classification, real numbers in regression).

Per-sample weights. Rescale C per sample. Higher weights force the classifier to put more emphasis on these points.

If X and y are not C-ordered and contiguous arrays of np.float64 and X is not a scipy.sparse.csr_matrix, X and/or y may be copied.

If X is a dense array, then the other methods will not support sparse matrices as input.

Get metadata routing of this object.

Please check User Guide on how the routing mechanism works.

A MetadataRequest encapsulating routing information.

Get parameters for this estimator.

If True, will return the parameters for this estimator and contained subobjects that are estimators.

Parameter names mapped to their values.

Perform classification on samples in X.

For an one-class model, +1 or -1 is returned.

For kernel="precomputed", the expected shape of X is (n_samples_test, n_samples_train).

Class labels for samples in X.

Compute log probabilities of possible outcomes for samples in X.

The model need to have probability information computed at training time: fit with

attribute probability set to True.

For `kernel="precomputed"`, the expected shape of `X` is `(n_samples_test, n_samples_train)`.

Returns the log-probabilities of the sample for each class in the model. The columns correspond to the classes in sorted order, as they appear in the attribute `classes_`.

The probability model is created using cross validation, so the results can be slightly different than those obtained by `predict`. Also, it will produce meaningless results on very small datasets.

Compute probabilities of possible outcomes for samples in `X`.

The model needs to have probability information computed at training time: fit with attribute probability set to True.

For `kernel="precomputed"`, the expected shape of `X` is `(n_samples_test, n_samples_train)`.

Returns the probability of the sample for each class in the model. The columns correspond to the classes in sorted order, as they appear in the attribute `classes_`.

The probability model is created using cross validation, so the results can be slightly different than those obtained by `predict`. Also, it will produce meaningless results on very small datasets.

Return accuracy on provided data and labels.

In multi-label classification, this is the subset accuracy which is a harsh metric since you require for each sample that each label set be correctly predicted.

Mean accuracy of `self.predict(X)` w.r.t. `y`.

Configure whether metadata should be requested to be passed to the fit method.

Note that this method is only relevant when this estimator is used as a sub-estimator within a meta-estimator and metadata routing is enabled with `enable_metadata_routing=True` (see `sklearn.set_config`). Please check the User Guide on how the routing mechanism works.

The options for each parameter are:

True: metadata is requested, and passed to fit if provided. The request is ignored if metadata is not provided.

False: metadata is not requested and the meta-estimator will not pass it to fit.

None: metadata is not requested, and the meta-estimator will raise an error if the user provides it.

str: metadata should be passed to the meta-estimator with this given alias instead of the original name.

The default (`sklearn.utils.metadata_routing.UNCHANGED`) retains the existing request. This allows you to change the request for some parameters and not others.

Added in version 1.3.

Metadata routing for `sample_weight` parameter in fit.

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as Pipeline). The latter have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.

Estimator parameters.

Configure whether metadata should be requested to be passed to the score method.

Note that this method is only relevant when this estimator is used as a sub-estimator within a meta-estimator and metadata routing is enabled with `enable_metadata_routing=True` (see `sklearn.set_config`). Please check the User Guide on how the routing mechanism works.

The options for each parameter are:

True: metadata is requested, and passed to score if provided. The request is ignored if metadata is not provided.

False: metadata is not requested and the meta-estimator will not pass it to score.

None: metadata is not requested, and the meta-estimator will raise an error if the user provides it.

str: metadata should be passed to the meta-estimator with this given alias instead of the original name.

The default (`sklearn.utils.metadata_routing.UNCHANGED`) retains the existing request. This allows you to change the request for some parameters and not others.

Added in version 1.3.

Metadata routing for `sample_weight` parameter in `score`.

Faces recognition example using eigenfaces and SVMs

Classifier comparison

Recognizing hand-written digits

Concatenating multiple feature extraction methods

Scalable learning with polynomial kernel approximation

Explicit feature map approximation for RBF kernels

Multilabel classification

ROC Curve with Visualization API

Custom refit strategy of a grid search with cross-validation

Statistical comparison of models using grid search