

tf.compat.v1.variable_axis_size_partitioner

Stay organized with collections Save and categorize content based on your preferences.

https://www.tensorflow.org/api_docs/python/tf/compat/v1/variable_axis_size_partitioner

Get a partitioner for VariableScope to keep shards below max_shard_bytes.

Function Signatures

```
tf.compat.v1.variable_axis_size_partitioner( max_shard_bytes,  
axis=0, bytes_per_string_element=16, max_shards=None )
```

Code Examples

```
tf.compat.v1.variable_axis_size_partitioner(  
max_shard_bytes, axis=0, bytes_per_string_element=16,  
max_shards=None  
)
```

```
tf.compat.v1.variable_axis_size_partitioner(  
max_shard_bytes, axis=0, bytes_per_string_element=16,  
max_shards=None  
)
```

Documentation

Get a partitioner for VariableScope to keep shards below max_shard_bytes.

This partitioner will shard a Variable along one axis, attempting to keep the maximum shard size below max_shard_bytes. In practice, this is not always possible when sharding along only one axis. When this happens, this axis is sharded as much as possible (i.e., every dimension becomes a separate shard).

If the partitioner hits the max_shards limit, then each shard may end up larger than max_shard_bytes. By default max_shards equals None and no limit on the number of shards is enforced.

One reasonable value for max_shard_bytes is $(64 \ll 20) - 1$, or almost 64MB, to keep below the protobuf byte limit.

Returns

Raises